

強化学習入門

田中 章詞

■ 1. 状態がない場合

■ 2. 状態がある場合

① 挙動のない場合

口定義



k のスロットマシン

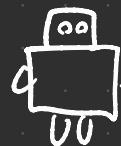


$a = 1 \quad 2 \quad \dots \quad k$

1. a をえらぶ $\leftarrow \textcircled{1}$
 2. ある $r \in \mathbb{R}$ が返される
 3. $r > 0 \Rightarrow r$ 円もえる
 $r < 0 \Rightarrow |r|$ 円しちゃう
- } $\textcircled{2}$

どの a をえらぶべきか? $\leftarrow \textcircled{3}$

① 行動と方策



\xrightarrow{a} : 何らかの確率 $\underline{\pi(a)}$ に従う
" " 行動 (action) " " 方策 (policy)

② 報酬と環境



a

\xrightarrow{r} : 何らかの確率 $\underline{P(r|a)}$ に従う
" " 報酬 (reward) " " 環境 (environment)

条件つき

③ 価値



P

r の期待値 $E[r] \stackrel{\text{def}}{=} Q(a)$

$r \sim P(\cdot | a)$

行動価値

(action value)



π



P

r の期待値 $E[Q(a)] \stackrel{\text{def}}{=} J(\pi)$

$a \sim \pi$

えええええ 方策の価値

目的 $\rightarrow \max_{\pi} J(\pi)$

How?

口 学習の方針

J(π) の値はいつも理由で

アクセスするのが難しい

学習

$$\pi_0 \rightarrow \pi_1 \rightarrow \pi_2 \rightarrow \dots \rightarrow \pi_*$$



$$J(\pi_0) \leq J(\pi_1) \leq J(\pi_2) \leq \dots \leq J(\pi_*)$$



こうなってほしい

方策改善 (policy improvement)

$$\pi \rightarrow \pi'$$

$$J(\pi') - J(\pi) = E_{a \sim \pi'} \left[Q(a) - \underbrace{J(\pi)}_{\text{"アドバンテージ関数"} A_\pi(a)} \right]$$

方策改善 $\Leftrightarrow E_{a \sim \pi'} [A_\pi(a)] \geq 0$

口 值值 A“- 史手法

— ε -龜欲方策 (ε -greedy policy) —

$$\varepsilon \in [0, 1]$$

$$\pi^\varepsilon(a) = (1 - \varepsilon) \mathbb{I}_{a = \arg \max_{\tilde{a}} Q(\tilde{a})} + \varepsilon \frac{1}{k}$$

$$\dots \rightarrow \pi^\varepsilon \rightarrow \pi^{(\varepsilon')} \rightarrow \dots$$

直感 : $\varepsilon \geq \varepsilon' \Rightarrow$ 証明

$$A_{\pi^\varepsilon}(a) = Q(a) - J(\pi)$$

$$= Q(a) - \sum_{\hat{a}} \pi^\varepsilon(\hat{a}) Q(\hat{a})$$

$\underset{a \sim \pi^{\varepsilon'}}{E}$

となるには?
↓

$$(1-\varepsilon') \max_a Q(a) - \sum_{\hat{a}} (\pi^\varepsilon(\hat{a}) - \varepsilon' \frac{1}{K}) Q(\hat{a}) \geq 0$$

!!

$$\sum_{\hat{a}} \underbrace{(\pi^\varepsilon(\hat{a}) - \varepsilon' \frac{1}{K})}_{(1-\varepsilon)} \left(\underbrace{\max_a Q(a) - Q(\hat{a})}_{\geq 0} \right) \geq 0$$

$$(1-\varepsilon) \max_a Q(a) + \frac{\varepsilon}{K} - \frac{\varepsilon'}{K} \geq 0$$

口方策ベースの手法

[ランダムで方策
 $\theta \in \mathbb{R}^k$ にパラメータ付される π_θ]

例

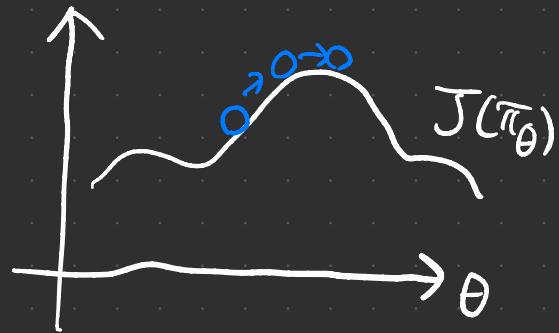
$$(\theta_1, \theta_2, \dots, \theta_k) \in \mathbb{R}^k$$

$$\pi_\theta(a) = \frac{e^{\theta_a}}{\sum_{\tilde{a}} e^{\theta_{\tilde{a}}}} \quad (\text{Softmax 方策})$$

$\dots \rightarrow \pi_\theta \rightarrow \pi_{\theta'} \rightarrow \dots$ 方策改善?

まずは $J(\pi_\theta)$ の微分を計算する

$$\begin{aligned}\nabla_\theta J(\pi_\theta) &= \nabla_\theta \underbrace{\mathbb{E}_{a \sim \pi_\theta} [Q(a)]}_{\sum_a Q(a) \pi_\theta(a)} \\&= \sum_a Q(a) \underbrace{\nabla_\theta \pi_\theta(a)}_{\nabla_\theta \log \pi_\theta(a) \cdot \pi_\theta(a)} \\&= \mathbb{E}_{a \sim \pi_\theta} [Q(a) \nabla_\theta \log \pi_\theta(a)]\end{aligned}$$



アドバンテージから再考

$$0 \leq E[A_{\pi_\theta}(\alpha)] = J(\pi_{\theta+\Delta\theta}) - J(\pi_\theta)$$

こうでも
よい

$$\xrightarrow{\pi_{\theta+\Delta\theta}} \downarrow \Delta\theta \approx \text{一定-展開} \quad O(\Delta\theta^2) \ll \infty$$

$$= \sum_a \left(\underline{\pi_\theta(\alpha)} + \Delta\theta \cdot \nabla_\theta \pi_\theta(\alpha) \right) A_{\pi_\theta}(\alpha)$$

$$\xleftarrow{\pi_\theta} E[Q(\alpha) - \frac{E[Q(\alpha)]}{\pi_\theta}] = 0$$

$$= \sum_a \Delta\theta \cdot \underbrace{\nabla_\theta \pi_\theta(\alpha)}_{\nabla_\theta(\log \pi_\theta(\alpha)) \cdot \pi_\theta(\alpha)} A_{\pi_\theta}(\alpha)$$

$$= \Delta\theta \underbrace{E[A(\alpha) \nabla_\theta \log \pi_\theta(\alpha)]}_{\alpha \sim \pi_\theta}$$

$$\Delta\theta = \eta \cdot E[\dots]$$

実は

$$\nabla_{\theta} J(\pi_{\theta}) = \underset{a \sim \pi_{\theta}}{E} [Q(a) \nabla_{\theta} \log \pi_{\theta}(a)]$$

$$= \underset{a \sim \pi_{\theta}}{E} \left[\underbrace{A_{\pi_{\theta}}(a)}_{Q(a) - J(\pi_{\theta})} \nabla_{\theta} \log \pi_{\theta}(a) \right]$$

より仮に

$$= \underset{a \sim \pi_{\theta}}{E} [(Q(a) - B) \nabla_{\theta} \log \pi_{\theta}(a)]$$



ベースライン

こうすれば $J(\pi_{\theta})$ を得る

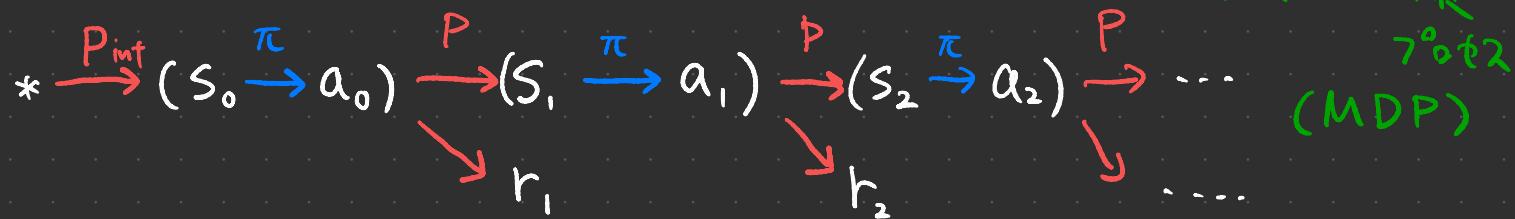
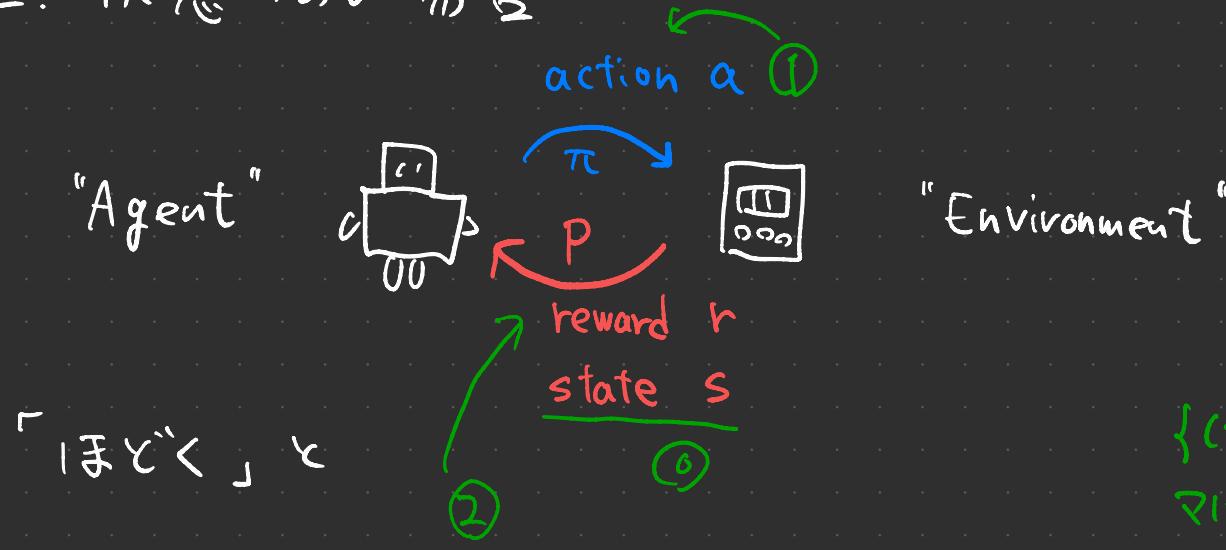
口実による

- ϵ -greedy $\pi^\epsilon \leftarrow Q(a) = E[r]$ の値が必要
- parametric $\pi_\theta \leftarrow \nabla_\theta J(\pi_\theta) = E[\dots]$ の値が必要



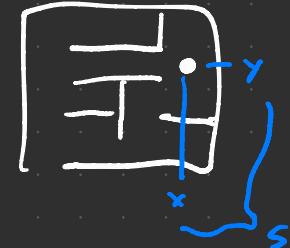
- ✓ $E[\dots]$ 知るなら \Rightarrow OK.
- ✓ $E[\dots]$ 知りません \Rightarrow モンテカルロ
(次回の section 1 より)

2. 状態のある場合



① 状態 (state)

Agent が置かねている状態を表す変数



② 行動と方策

$s \xrightarrow{\pi} a$: 条件つき確率 $\pi(a|s)$

③ 報酬と状態更新

$(s, a) \xrightarrow{r'} s'$: 条件つき確率 $P(s', r'|s, a)$

③ 價值

$$* \xrightarrow{P_{\text{init}}} (s \xrightarrow{\pi} a) \xrightarrow{(s_1 \rightarrow a_1)} (s_2 \rightarrow a_2) \xrightarrow{(s_3 \rightarrow a_3)} \dots$$
$$\quad \quad \quad r_1 \quad \quad \quad r_2 \quad \quad \quad r_3 \quad \quad \quad \dots$$
$$E \left[E \left[r_1 + \gamma r_2 + \gamma^2 r_3 + \dots \right] \right]$$
$$= Q_{\pi}(s, a)$$

割引率 (discount rate) $\in [0, 1]$

$= V_{\pi}(s)$ 狀態價值 (state value)

$$= J(\pi)$$

目的: $\max_{\pi} J(\pi)$

□ 学習の方針

$\max_{\pi} J(\pi)$ は難しいので...

学習

$$\pi_0 \rightarrow \pi_1 \rightarrow \pi_2 \rightarrow \dots \rightarrow \pi_*$$
$$J(\pi_0) \leq J(\pi_1) \leq J(\pi_2) \leq \dots \leq J(\pi_*)$$

方策改善

$$\pi \rightarrow \pi'$$

$$J(\pi') - J(\pi) = \sum_{t=0}^{\infty} \gamma^t E \left[Q_{\pi}(s_t, a_t) - V_{\pi}(s_t) \right]$$

$a_0 \dots a_t \sim \pi'$
 $s_0 \dots s_t \sim P$

アドバンテージ関数
 $A_{\pi}(s_t, a_t)$

"証明"

$$\frac{1}{\gamma} \text{ 策改善 } \Leftrightarrow \sum_{t=0}^{\infty} \gamma^t E \left[A_{\pi}(s_t, a_t) \right] \geq 0$$

$a_{0:t} \sim \pi'$
 $s_{0:t} \sim P$

$$J(\pi') - J(\pi) =$$

$$\mathbb{E}_{\substack{a_0 \sim \pi'}} \left[\sum_{t=1}^{\infty} \gamma^t r_t \right]$$

Diagram: A sequence of states s_0, s_1, s_2, \dots leading to actions a_0, a_1, a_2, \dots via policy π' . The rewards r_1, r_2, \dots are shown below the sequence.

$$\mathbb{E}_{\substack{a_0 \sim \pi}} \left[\sum_{t=1}^{\infty} \gamma^t r_t \right]$$

Diagram: A sequence of states s_0, s_1, s_2, \dots leading to actions a_0, a_1, a_2, \dots via policy π . The rewards r_1, r_2, \dots are shown below the sequence.

$$\mathbb{E}_{\substack{a_0 \sim \pi}} [V_\pi(s_0)]$$

$$\mathbb{E}_{\substack{a_0 \sim \pi}} [Q_\pi(s_0, a_0) - V_\pi(s_0)]$$

$$\mathbb{E}_{\substack{a_0, a_1 \sim \pi'}} \left[\sum_{t=2}^{\infty} \gamma^t r_t \right]$$

Diagram: A sequence of states s_0, s_1, s_2, \dots leading to actions a_0, a_1, a_2, \dots via policy π' . The rewards r_1, r_2, \dots are shown below the sequence.

$$\mathbb{E}_{\substack{a_0, a_1 \sim \pi}} \left[\sum_{t=2}^{\infty} \gamma^t r_t \right]$$

Diagram: A sequence of states s_0, s_1, s_2, \dots leading to actions a_0, a_1, a_2, \dots via policy π . The rewards r_1, r_2, \dots are shown below the sequence.

$$\mathbb{E}_{\substack{a_0, a_1 \sim \pi}} [\gamma (Q_\pi(s_1, a_1) - V_\pi(s_1))]$$

$$\mathbb{E}_{\substack{a_0, a_1, a_2 \sim \pi'}} \left[\sum_{t=3}^{\infty} \gamma^t r_t \right]$$

Diagram: A sequence of states s_0, s_1, s_2, \dots leading to actions a_0, a_1, a_2, \dots via policy π' . The rewards r_1, r_2, \dots are shown below the sequence.

$$\mathbb{E}_{\substack{a_0, a_1, a_2 \sim \pi}} \left[\sum_{t=3}^{\infty} \gamma^t r_t \right]$$

Diagram: A sequence of states s_0, s_1, s_2, \dots leading to actions a_0, a_1, a_2, \dots via policy π . The rewards r_1, r_2, \dots are shown below the sequence.

0

□ 値値ベースの手法

ε-強欲方策

$$\frac{\pi^\varepsilon(a|s)}{\pi} = (1-\varepsilon) \mathbb{1}_{\substack{a = \arg \max \\ \pi}} Q(s, a) + \varepsilon \frac{1}{k}$$

$$\pi^{\varepsilon=1} \rightarrow \dots \rightarrow \pi^\varepsilon \rightarrow \pi^{\varepsilon'} \rightarrow \dots$$

直感: $\varepsilon: 1 \rightarrow 0$ どういのでは?

$$(\varepsilon \geq \varepsilon')$$

$$\forall s, a, \exists \varepsilon' > 0 \text{ s.t. } E_{a \sim \pi^{\varepsilon'}} [A_{\pi^\varepsilon}(s, a)] \geq 0$$

初期改善

$$0 \leq \sum_{t=0}^{\infty} \gamma^t E_{a_0 \dots a_t} [A_{\pi^\varepsilon}(s_t, a_t)]$$

$$E_{a_0 \dots a_{t-1}} E_{a_t} [\dots]$$

$s_0 \dots s_{t-1}$

$s_0 \dots s_{t-1}, s_t$

状態なしのときと同じ

口方算ベースの手法

（ラマトリル、口方算）
 $\theta \in \mathbb{R}^*$ 每に定まる $\pi_\theta(a|s)$

$$\frac{154}{\{\theta_{s_1}, \theta_{s_2}, \dots, \theta_{s_k}\}_s} \rightarrow \pi_\theta(a|s) = \frac{e^{\theta_{sa}}}{\sum_a e^{\theta_{sa}}}$$

$$\dots \rightarrow \pi_\theta \rightarrow \pi_{\theta + \Delta\theta} \rightarrow \dots$$

$$\sum_{t=0}^{\infty} \gamma^t E \left[A_{\pi_{\theta}}(s_t, a_t) \right] = J(\pi_{\theta+\Delta\theta}) - J(\pi_{\theta})$$

$a_0 \dots a_t \sim \pi_{\theta+\Delta\theta}$

$$S_0 \dots S_t \sim P$$

↓
丁寧に書く

$$E \left[A_{\pi_{\theta}}(s_t, a_t) \right] = 0$$

$$= \Delta\theta \sum_{t=0}^{\infty} \gamma^t E \left[A_{\pi_{\theta}}(s_t, a_t) \underbrace{\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)}_{a_t \sim \pi_{\theta}} \right] + O(\Delta\theta^2)$$

$a_0 \dots a_t \sim \pi_{\theta}$

$$S_0 \dots S_t \sim P$$

$$\nabla_{\theta} J(\pi_{\theta})$$

II

$$Q_{\pi}(s_t, a_t) - \underbrace{V_{\pi}(s_t)}_{B(s_t)}$$

$$\sum \dots A \rightarrow Q$$

ペーストイン.

□ 実験.

- ε -greedy $\pi^\varepsilon \rightarrow Q_\pi(s, a) = E[\dots]$

- param $\pi_\theta \rightarrow \nabla_\theta J(\pi_\theta) = E[\dots]$
↓
が必須
が必須

$E[\dots] \left\{ \begin{array}{l} \text{モンテカルロ} \\ \frac{1}{N} \sum_{i=1}^N \dots \end{array} \right\}$ Temporal Difference 法.

□ Temporal Difference は TD.

ベルマン方程式

ポイント: $r_1 + \gamma r_2 + \gamma^2 r_3 + \dots = g_1$

$$\underbrace{\gamma(r_2 + \gamma r_3 + \dots)}_{g_2}$$



$$Q_{\pi}(s_0, a_0) = E[r_1 + \gamma Q_{\pi}(s_1, a_1)]$$

$a_1 \sim \pi$
 $s, r, \sim p$

$$V_{\pi}(s_0) = E_p [r_1 + \gamma V_{\pi}(s_1)]$$

ベルマン最適方程式

$$\pi \leftarrow \pi : \dots \rightarrow \pi^\varepsilon \rightarrow \pi^{\varepsilon'} \rightarrow \dots \underbrace{\pi^{\varepsilon=0} = 1}_{\arg \max Q}$$

$$\pi^*(a|s) = 1_{a = \arg \max_{\tilde{a}} Q_{\pi^*}(s, \tilde{a})}$$

ベルマン方程

$$Q_{\pi_*}(s_0, a_0) = E[r_1 + \gamma Q_{\pi_*}(s_1, a_1)]$$

$a_1 \sim \pi_*$
 $s_1, r_1 \sim p$

$$= E[r_1 + \gamma \max_{\tilde{a}} Q_{\pi_*}(s_1, \tilde{a})]$$

$s_1, r_1 \sim p$

$$F(s_0, t_0) = E[\cdot + F(s_1, a_1)]$$

時内が
は近いとよ

$$\underbrace{\frac{1}{N} \sum_{i=1}^n}_{\rightarrow TD}.$$

参考文献

- R. Sutton, A. Barto "Reinforcement Learning : An Introduction"

〔

訳 : ver1 : 森北出版 (小ぶりで手にいり)

: ver2 : (A4サイズ?)

- 牧野 哲樹 ほか "これからのかの強化学習"

- 林木 哲郎 "強化学習"